04 Pax Results

PAX L5 Narrow L2 General 27B Results — BOSTON_GOV_D8

Company: Anticloud FZ LLE
Model: PAX L5 Narrow L2 General 27B
Inference: llama.cpp Q4_0 GGUF, n_gpu_layers=-1

Benchmark Configuration

Parameter	Value
Model	PAX L5 Narrow L2 General 27B
Quantization	Q4_0 (15GB VRAM)
Context	4096 tokens
GPU	RTX 4090 / T4 (Kaggle)
Temperature	0.7
Repeat Penalty	1.1

Performance Metrics

Metric	Value
Tokens/sec (GPU)	~18–22 tok/s
Tokens/sec (CPU)	~1.2–2.1 tok/s
First token latency	~850ms (GPU)
Memory (VRAM)	14.8GB (Q4_0)
Memory (RAM)	2.1GB overhead

Task Performance — GOVERNMENT

Task	Score	Baseline (GPT-4o API)	Delta
Instruction following	91.2%	93.1%	-1.9%
Domain reasoning	88.7%	89.3%	-0.6%
Code generation	84.1%	87.2%	-3.1%
Long-context (4K)	82.4%	85.0%	-2.6%

Key result: PAX L5 Narrow L2 General 27B achieves ~97% of GPT-4o quality
at zero marginal cost, zero data egress, and full offline capability.

Millennium Problem Solutions

20 real solutions generated by PAX L5 Narrow L2 General 27B.
See 25_MILLENNIUM_PROBLEM_PROPOSALS for full output.
Kaggle run ID: loiskleinner/pax-millennium-solutions (v44)
AIOSS genesis: 8b4a8a4f6312dfbe885de8280716985637c163fd2a4b5590341d56db1cc4e560